Storage of Databases

- Need data to be stored “permanently” or *persistently* for long periods of time

File Organizations

- File organization refers to the way records are stored in terms of blocks and the way blocks are placed on the storage medium.
- Types of organizations
  - Unsorted
  - Sorted
  - Hashing

Records

- A file is a *sequence* of records
- Records could be either *fixed-length* or *variable-length*
- Records comprise of a sequence of *fields* (column, attribute)

Blocks

- A block is **the largest contiguous amount of disk space that can be allocated to a file** and is therefore the largest amount of data that can be accessed in a single I/O operation.

Blocking Factor

The number of records that are stored in a block is called the "blocking factor". Blocking factor is constant across blocks if record length is fixed, or variable otherwise.

If B is block size and R is record size, then blocking factor is:

$$bfr = \lfloor B/R \rfloor$$

Since R may not exactly divide B, there could be some left-over space in each block equal to:

$$B - (bfr * R) \quad \text{bytes.}$$

Spanned and Unspanned Records

When extra space in blocks are left unused, the record organization is said to be “unspanned”.

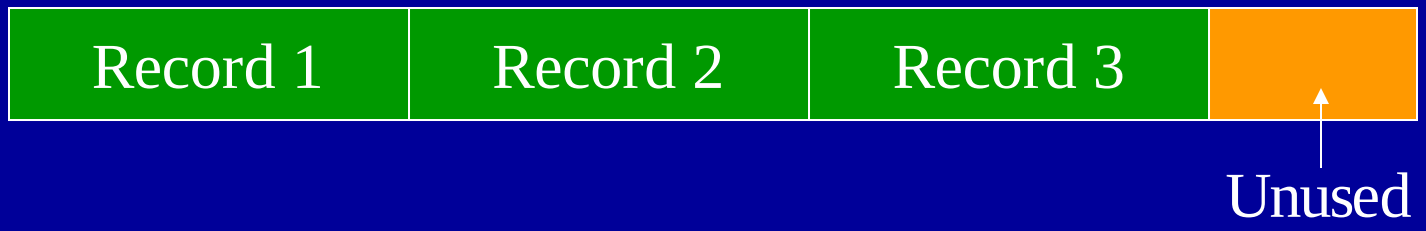

Spanned and Unspanned Records

In “spanned” record storage, records can be split so that the “span” across blocks.

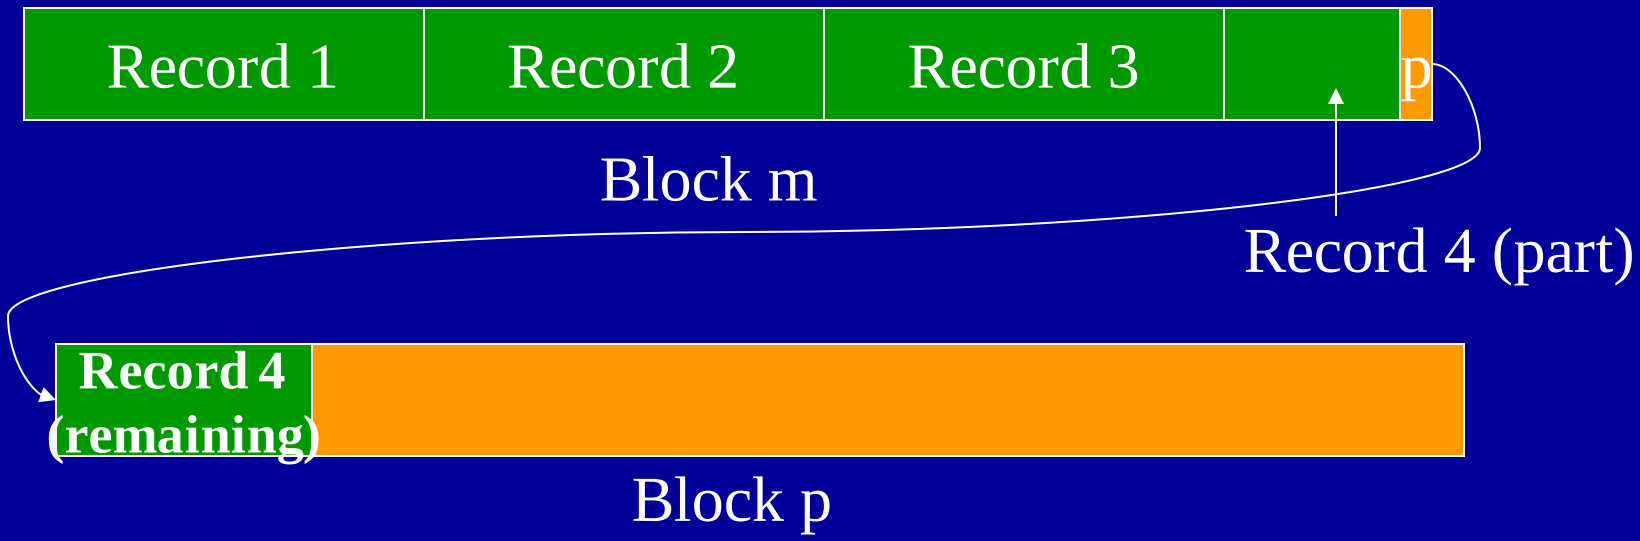

Spanned and Unspanned Records

When record size is greater than block size (i.e. R > B), use of "spanned" record storage is compulsory.

Unordered File Organization

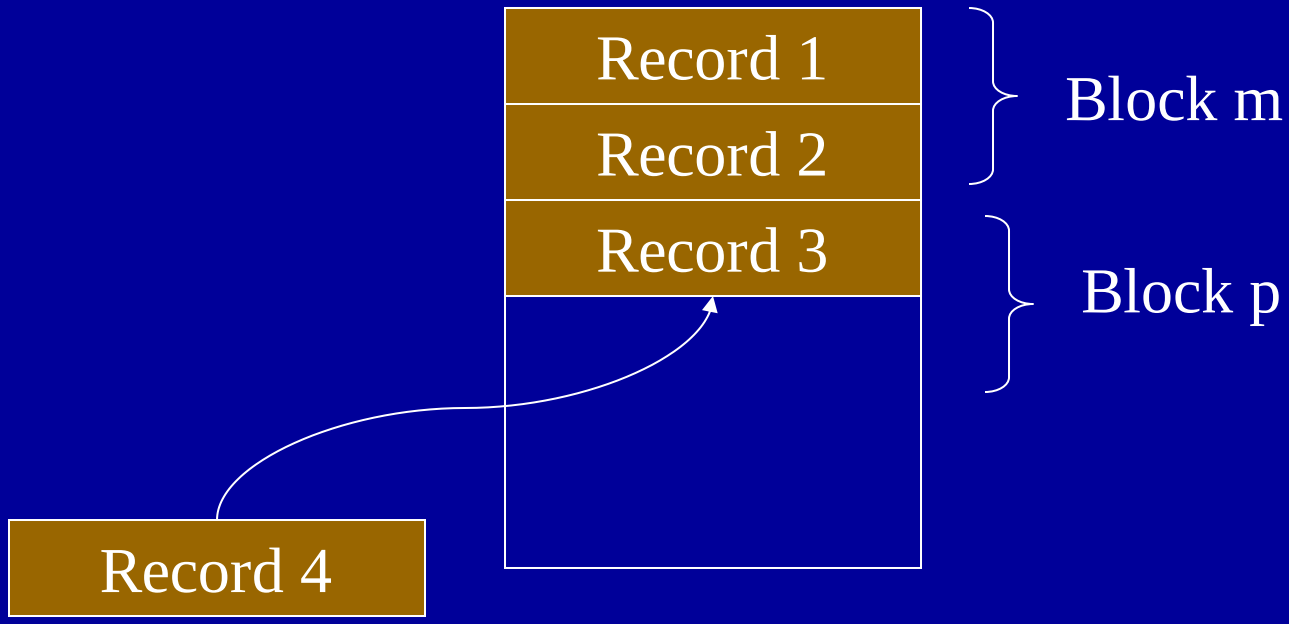

Also called a “pile file”…

Pile Files

- Simplest file organization
- Records are inserted in the order of their arrival
- Insertion: very easy

Record Deletion in Pile Files

Deletion by fragmentation.
Inefficient in terms of space usage.

Sorted Files

- File organization where records are sorted based on the value of some field called the ordering field
- Ordering field should be a key field (unique for each record) and belong to an ordinal domain (where values can be compared)
- Insertion and Deletion: both expensive
- Updation may involve physical migration of record, if ordering field is modified
- Searching: Binary search based on ordering field efficient

Sorted Files

(Sorted)

Sorted Files

- More efficient than pile-files for key-based searches
- A better choice when database is mostly read-only and queries are mostly key-based retrievals

Hashing Techniques

- Provide very fast access to records on certain search conditions
- Search condition should be an equality condition on a “key” field
- Uses a “hashing function” to map keys onto “buckets” hosting records

Hashing

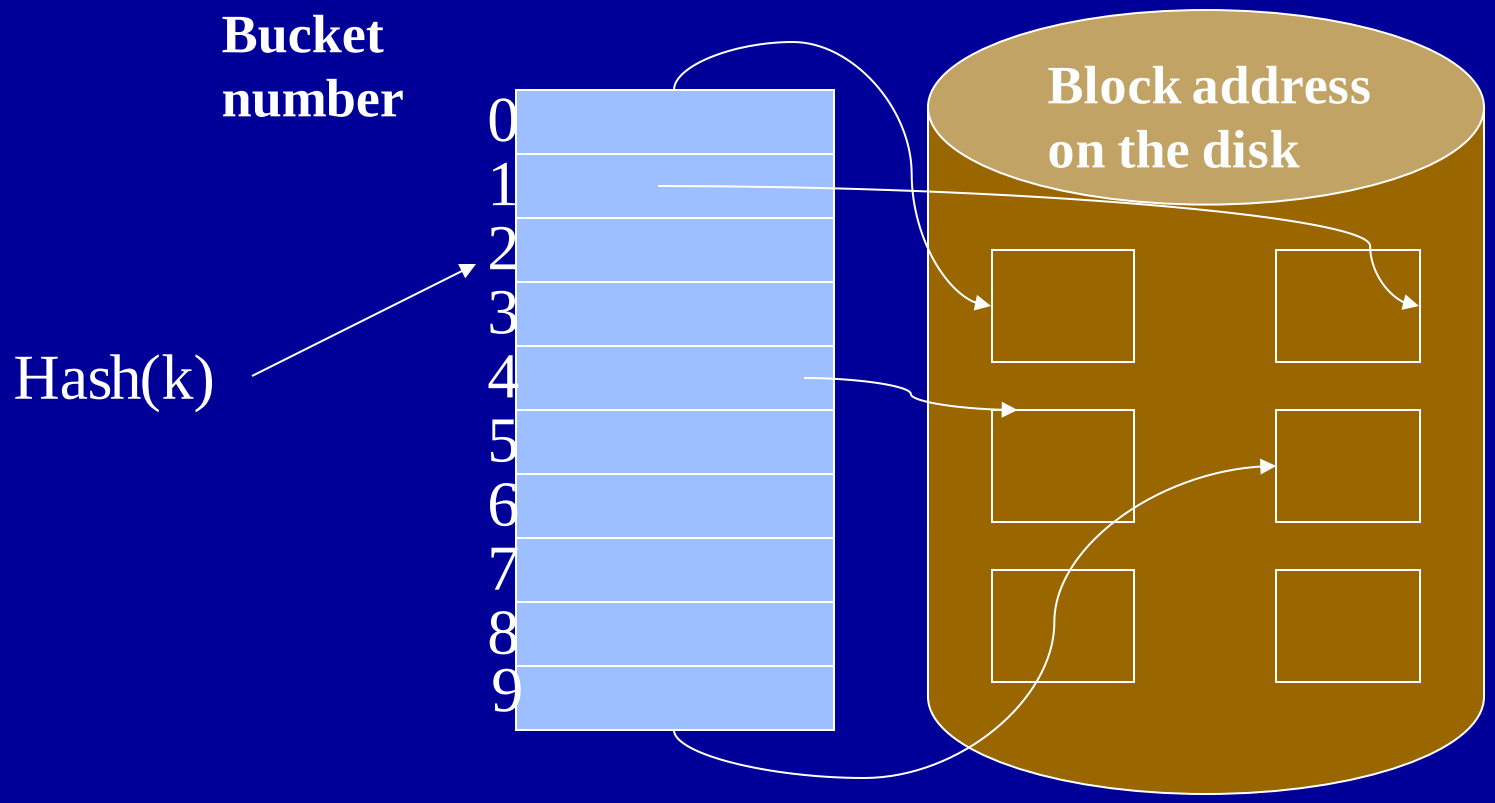

Hashing

- Uses two levels of indirection: hashing to buckets and searching within buckets
- A bucket is a disk block or a set of contiguous blocks
- A bucket can hold more than one record
- Sequential search within buckets

Overflows

Occurs when a bucket receives more records than it can hold. Overflow management techniques:

- Open addressing
  - Use the next available bucket.
- Chaining
  - Pointers to extra buckets called overflow buckets.
- Rehashing
  - Use a second hashing function to see if it can hash without overflow.

Dynamic Hashing

- Hashing where the number of buckets are fixed are called *static* hashing
- Inefficient when data set is skewed. Few overflowing buckets and large number of nearly empty buckets
- Extensible or Dynamic hashing: number of buckets change dynamically in accordance with data requirements

Dynamic Hashing

- No. of buckets may grow or shrink depending on additions and deletions
- Overall strategy:
  1. start with a single bucket.
  2. Once full, split the bucket into two.
  3. Redistribute records between the two split buckets using some method ensuring near-uniform distribution

Dynamic Hashing

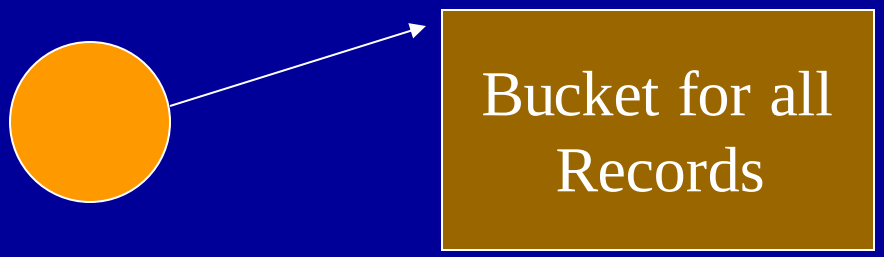

Assume keys to be made of binary strings…

Dynamic Hashing

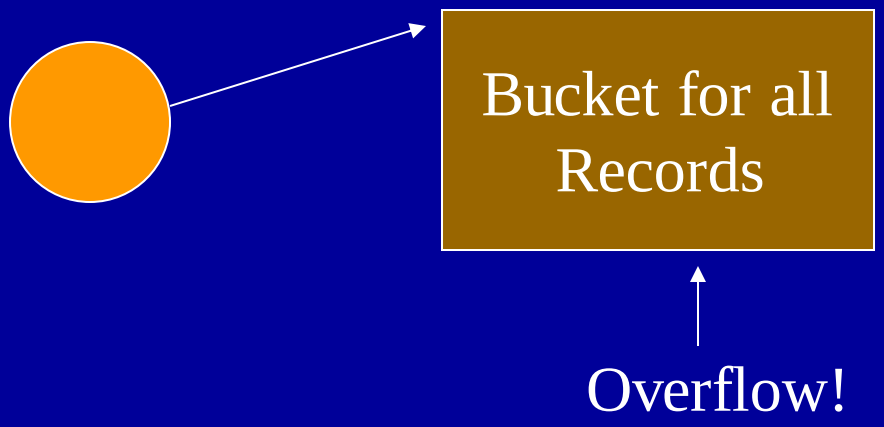

Dynamic Hashing

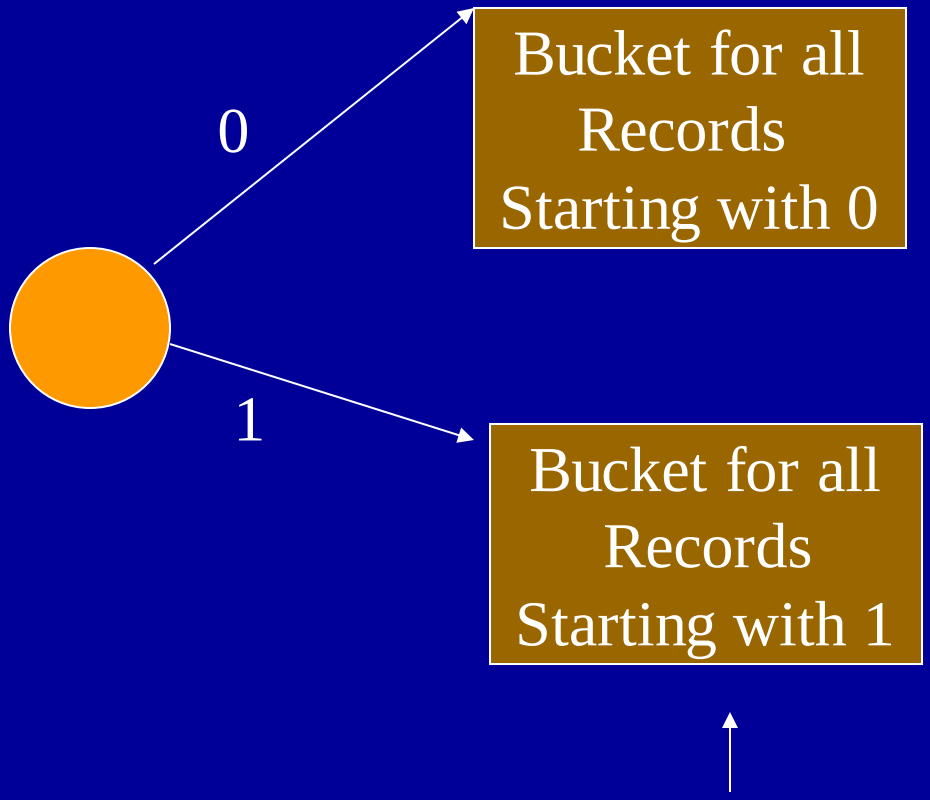

Dynamic Hashing

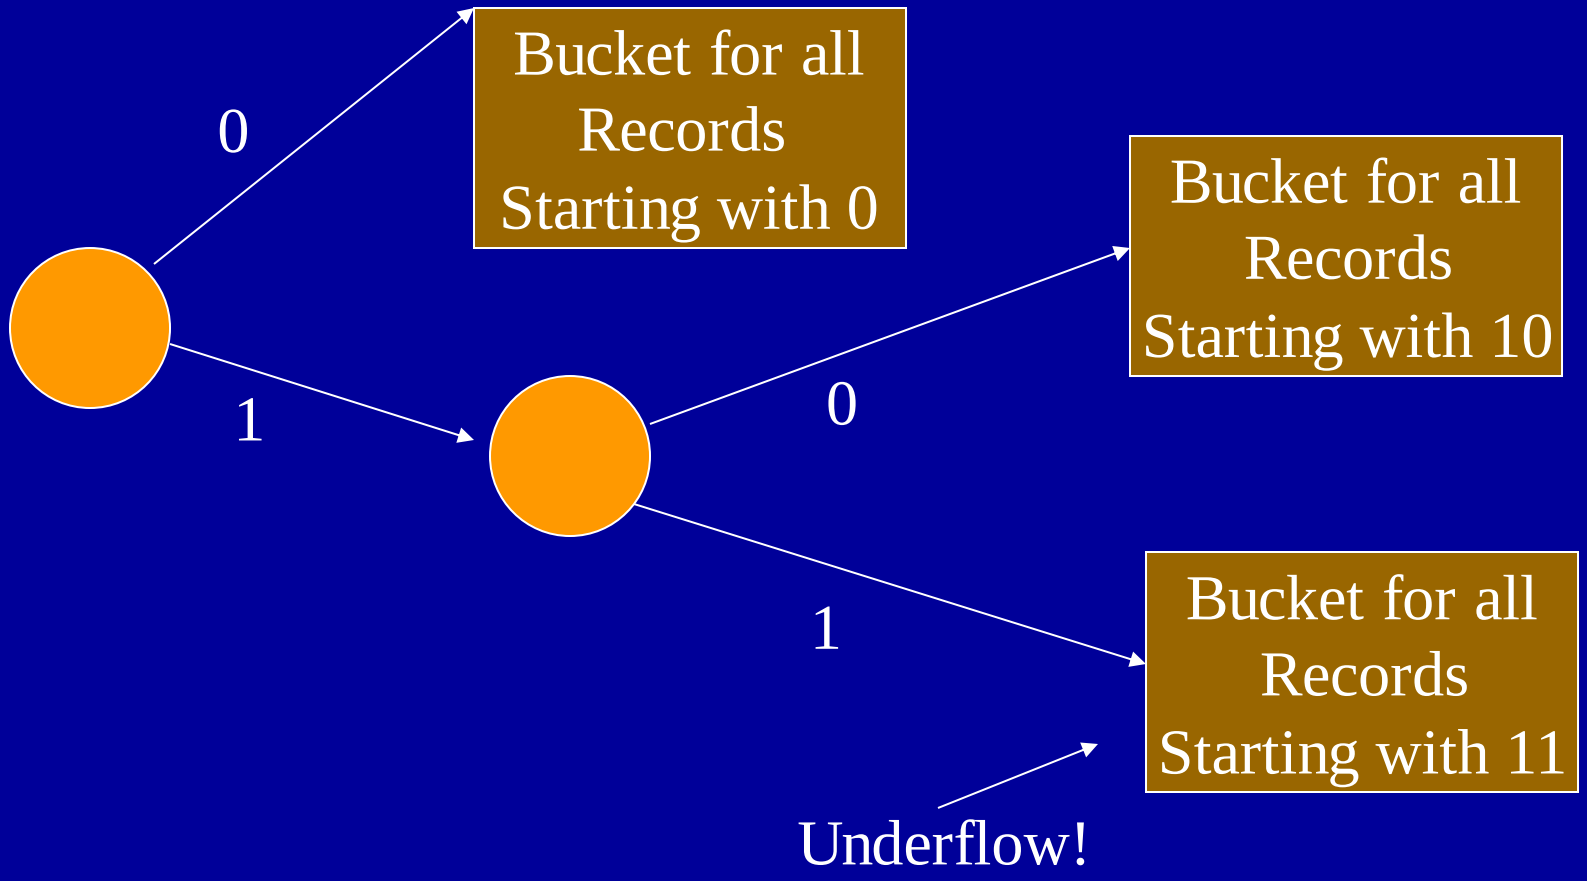

Dynamic Hashing

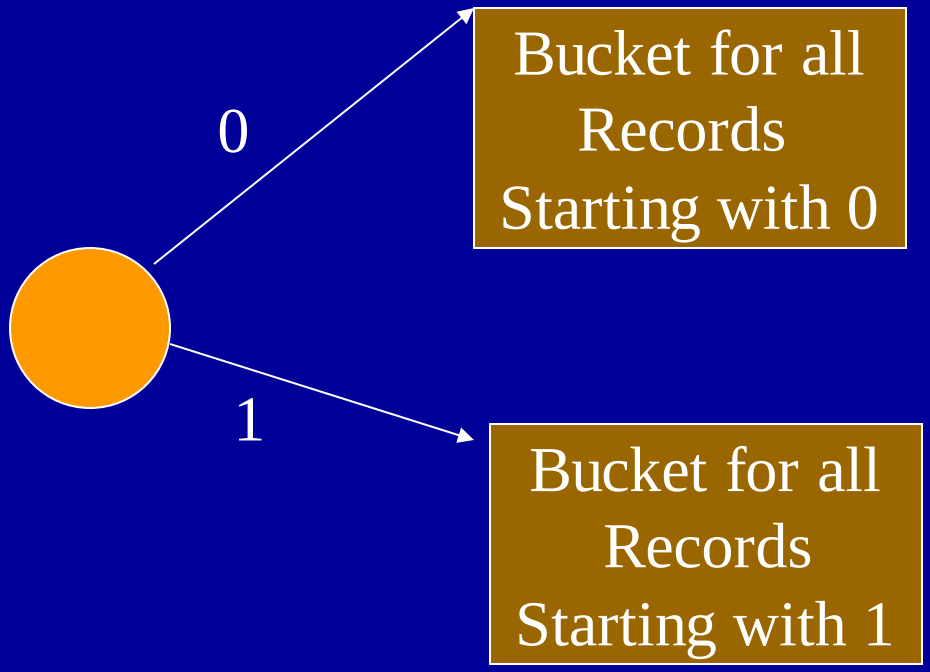

Summary

- Storage Media and their characteristics
- Records, Blocks and Files
- Spanning and Unspanning records
- Pile Files
- Sorted Files
- Hash Files
  - Static versus dynamic hashing