

# Physical Design


- Design so far
  - E/R modelling helps find the requirements of a database
  - Normalisation helps to refine a design by removing data redundancy
- Physical design
  - Concerned with storing and accessing the data
  - How to deal with media failures
  - How to access information efficiently

# RAID Arrays

- RAID - redundant array of independent (inexpensive) disks
  - Storing information across more than one physical disk
  - Speed - can access more than one disk
  - Robustness - if one disk fails it is OK
- RAID techniques
  - Mirroring - multiple copies of a file are stored on separate disks
  - Striping - parts of a file are stored on each disk
  - Different levels (RAID 0, RAID 1...)


# RAID Level 0

- Files are split across several disks
  - For a system with  $n$  disks, each file is split into  $n$  parts, one part stored on each disk
  - Improves speed, but no redundancy



# RAID Level 1

- As RAID 0 but with redundancy
  - Files are split over multiple disks
  - Each disk is mirrored
  - For  $n$  disks, split files into  $n/2$  parts, each stored on 2 disks
  - Improves speed, has redundancy, but needs lots of disks



# Parity

- Parity - for a set of data in binary form we count the number of 1s for each bit across the data
- If this is even the parity is 0, if odd then it is 1

1 0 1 1 0 0 1 1

0 0 1 1 0 0 1 1

1 0 1 0 1 0 0 1

0 1 1 0 1 1 1 0

---

0 1 0 0 0 1 1 1


# Recovery With Parity

- If one of our pieces of data is lost we can recover it
  - Just compute it as the parity of the remaining data and our original parity information

$$\begin{array}{cccccccc} 1 & 0 & 1 & 1 & 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 & 0 & 1 & 1 \\ & & & & & & & \\ 0 & 1 & 1 & 0 & 1 & 1 & 1 & 0 \\ \hline 0 & 1 & 0 & 0 & 0 & 1 & 1 & 1 \end{array}$$

# RAID Level 3

- Data is striped over disks, and a parity disk for redundancy
  - For  $n$  disks, we split the data in  $n-1$  parts
  - Each part is stored on a disk
  - The final disk stores parity information



# Other RAID Issues

- Other RAID levels consider
  - How to split data between disks
  - Whether to store parity information on one disk, or spread across several
  - How to deal with multiple disk failures



# Indexes (discussed)

# Query Processing

- Once a database is designed and made we can query it
  - A query language (such as SQL) is used to do this
  - The query goes through several stages to be executed
- Three main stages
  - Parsing and translation - the query is put into an internal form
  - Optimisation - changes are made for efficiency
  - Evaluation - the optimised query is applied to the DB

# Parsing and Translation

- SQL is a good language for people
  - It is quite high level
  - It is non-procedural
- Given an SQL statement we want to find an equivalent relational algebra expression
- This expression may be represented as a tree - the query tree

# Some Relational Operators

- Product  $\times$ 
  - Product finds all the combinations of one tuple from each of two relations
  - $R1 \times R2$  is equivalent to  
**SELECT \***  
**FROM R1, R2**
- Selection  $\sigma$ 
  - Selection finds all those rows where some condition is true
  - $\sigma_{\text{cond}} R$  is equivalent to  
**SELECT \***  
**FROM R**  
**WHERE <cond>**

# Some Relational Operators

- Projection  $\pi$ 
  - Projection chooses a set of attributes from a relation, removing any others

- $\pi_{A1,A2,\dots} R$  is equivalent to  

```
SELECT
    A1, A2, ...
FROM R
```

- Projection, selection and product are enough to express queries of the form

```
SELECT <cols>
FROM <table>
WHERE <cond>
```

# SQL → Relational Algebra

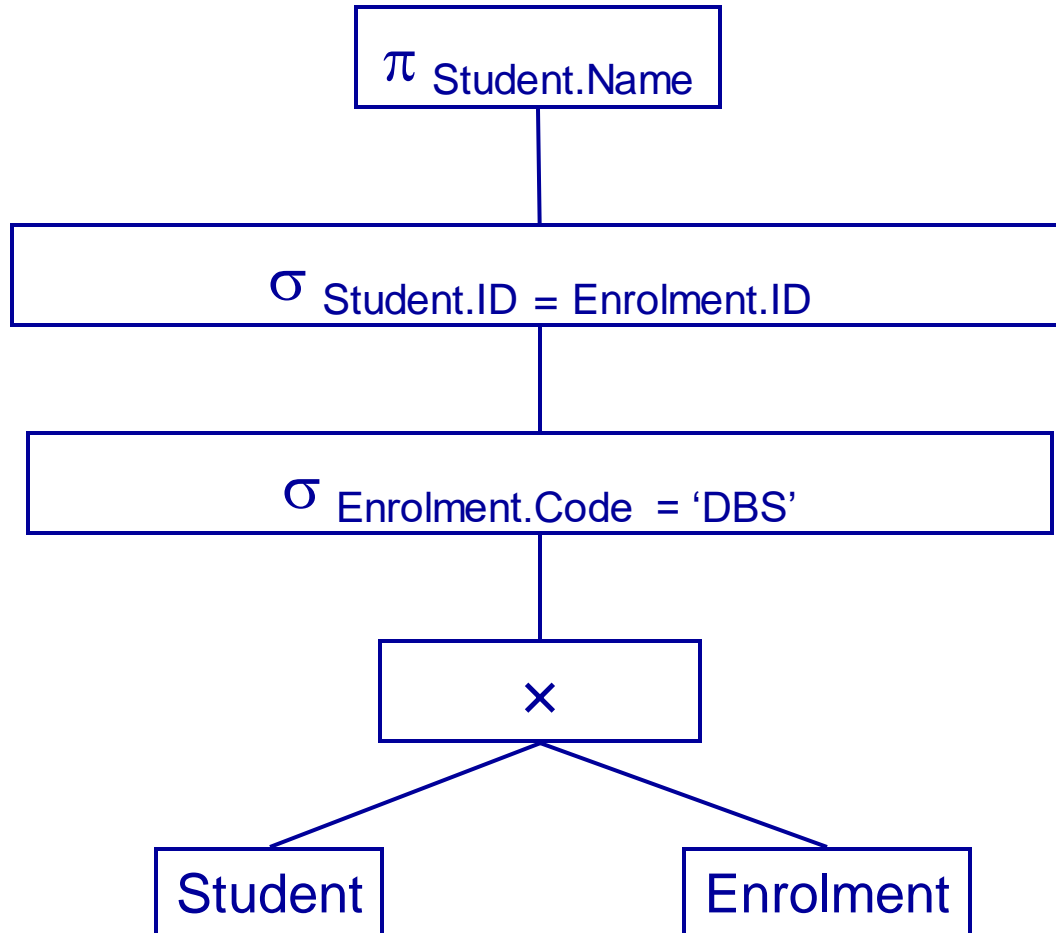
- SQL statement

```
SELECT Student.Name
FROM Student,
      Enrolment
WHERE
      Student.ID =
      Enrolment.ID
AND
      Enrolment.Code =
      'DBS'
```

- Relational Algebra

- Take the product of Student and Enrolment
- select tuples where the IDs are the same and the Code is DBS
- project over Student.Name

# Query Tree



# Optimisation

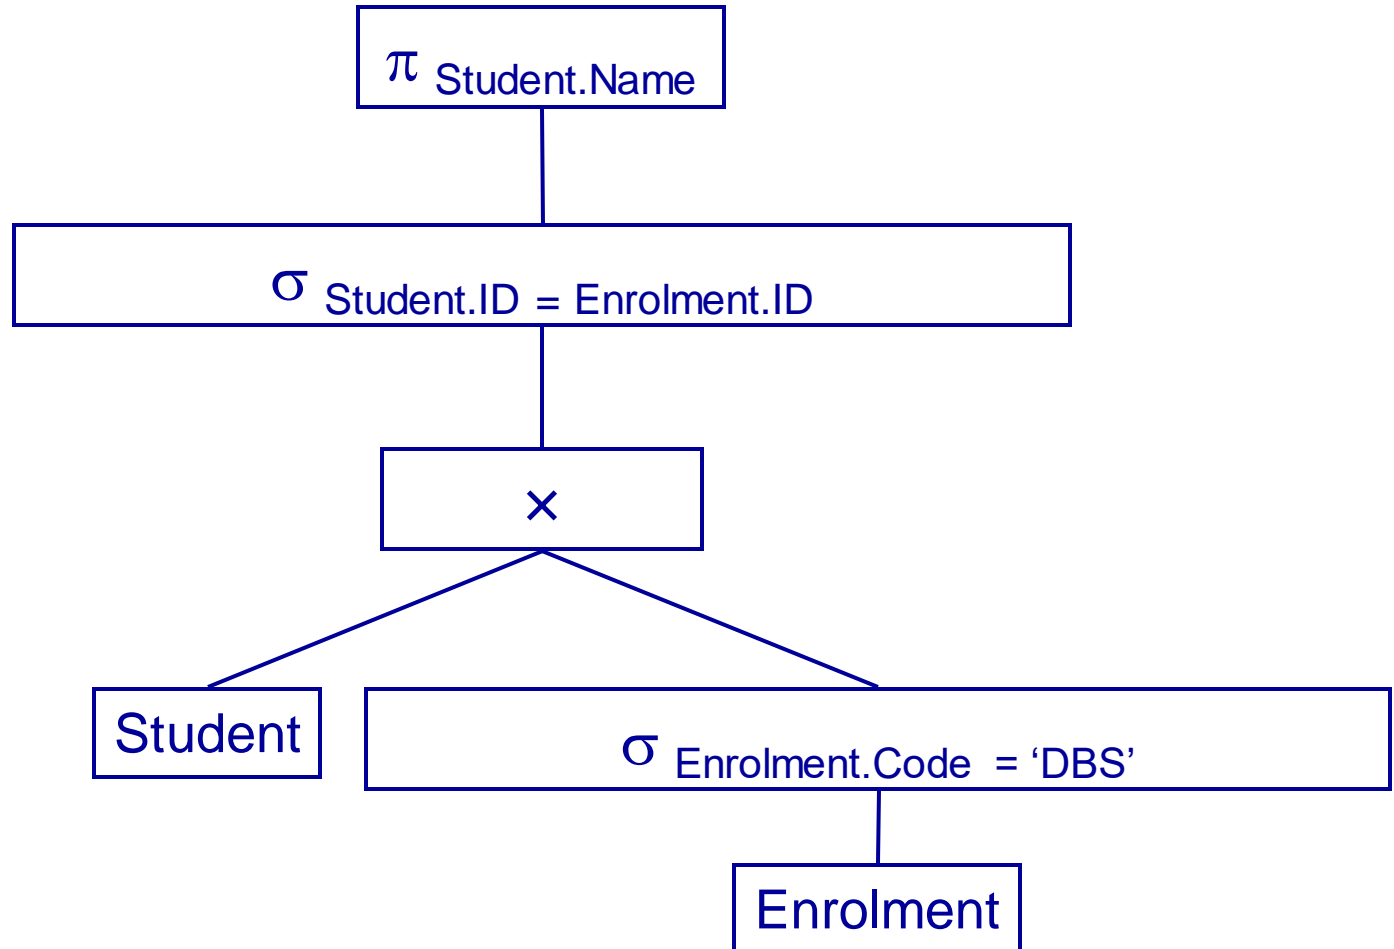
- There are often many ways to express the same query
- Some of these will be more efficient than others
- Need to find a good version
- Many ways to optimise queries
  - Changing the query tree to an equivalent but more efficient one
  - Choosing efficient implementations of each operator



# Optimisation Example

- In our query tree before we have the steps
  - Take the product of Student and Enrolment
  - Then select those entries where the Enrolment.Code equals 'DBS'
- This is equivalent to
  - selecting those Enrolment entries with Code = 'DBS'
  - Then taking the product of the result of the selection operator with Student

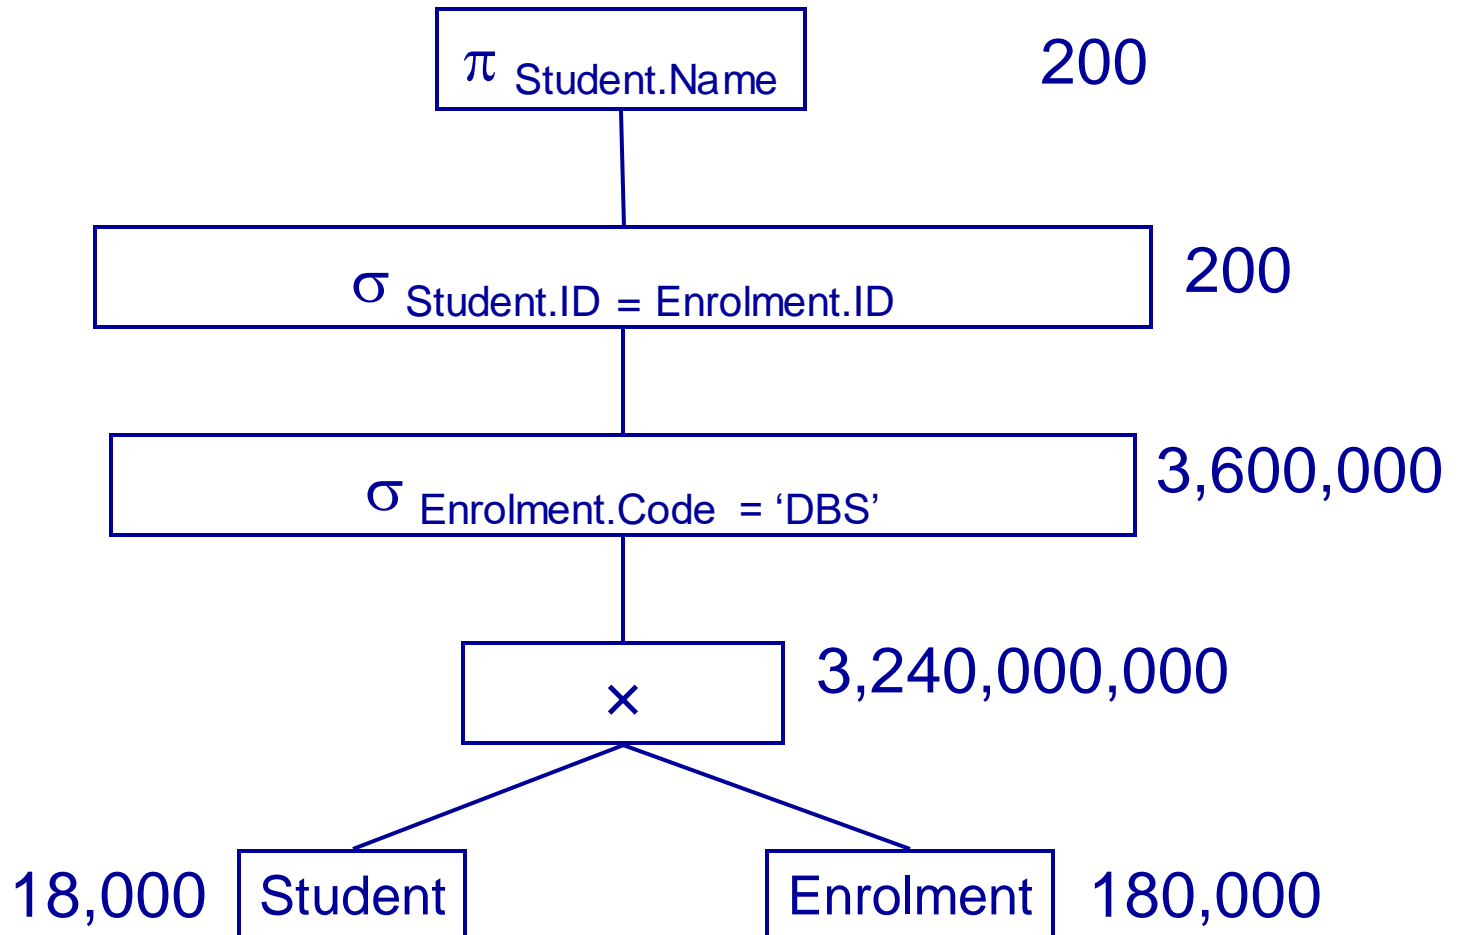
# Optimised Query Tree



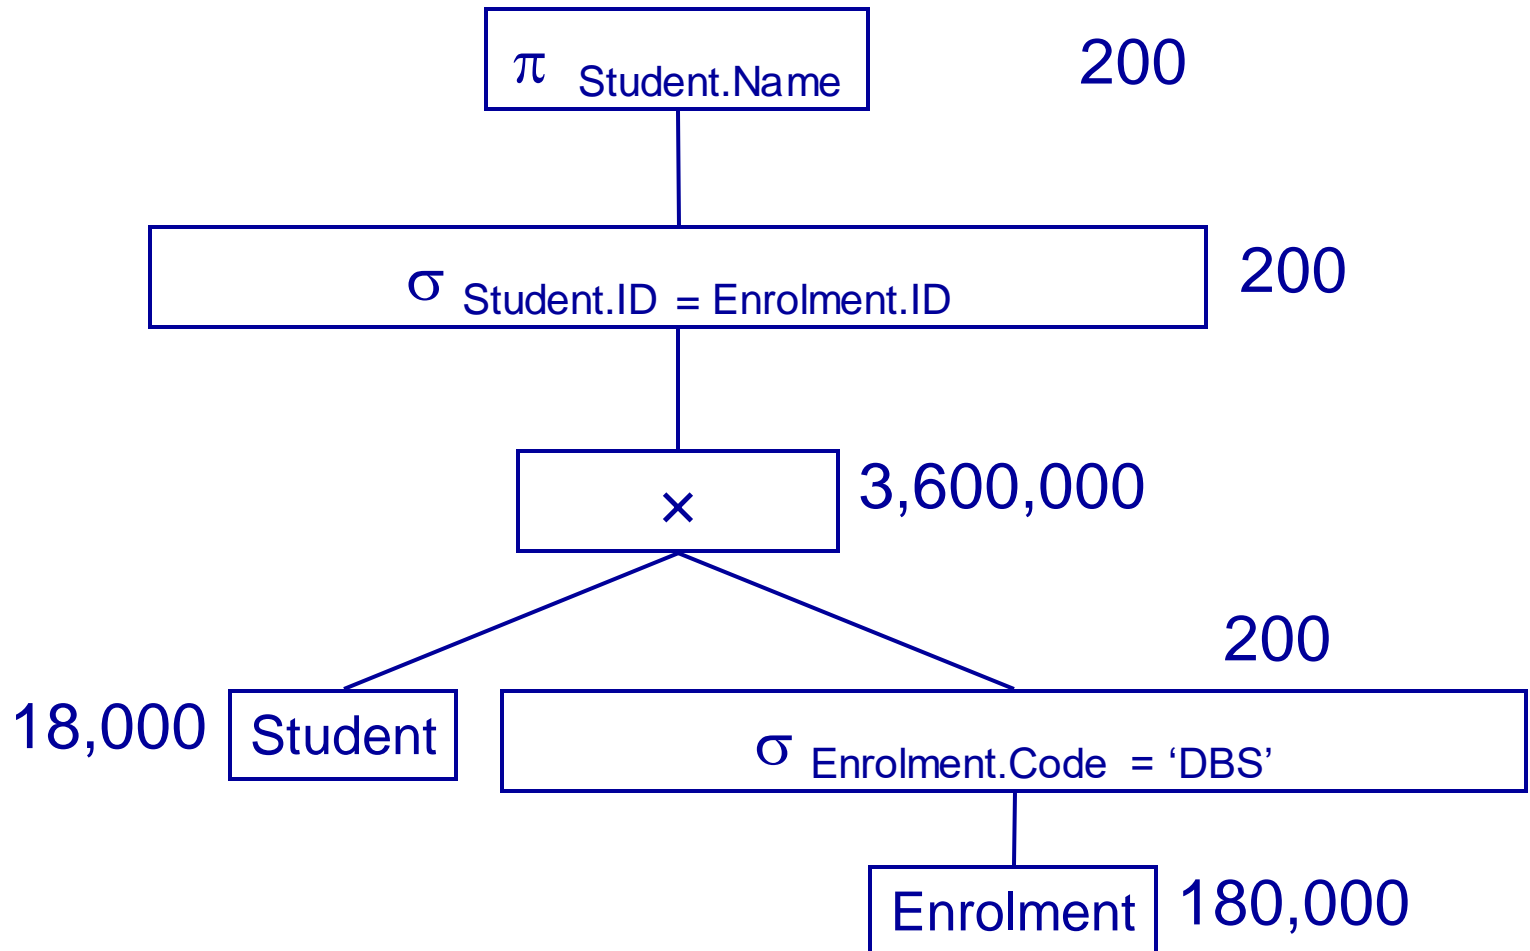
# Optimisation Example

- To see the benefit of this, consider the following statistics
  - Nottingham has around 18,000 full time students
  - Each student is enrolled in at about 10 modules
  - Only 200 take DBS
- From these statistics we can compute the sizes of the relations produced by each operator in our query trees

# Original Query Tree



# Optimised Query Tree



# Optimisation Example

- The original query tree produces an intermediate result with 3,240,000,000 entries
- The optimised version at worst has 3,600,000
- A big improvement!
- There is much more to optimisation
  - In the example, the product and the second selection can be combined and implemented efficiently to avoid generating all Student-Enrolment combinations

# Optimisation Example

- If we have an index on Student.ID we can find a student from their ID with a binary search
- For 18,000 students, this will take at most 15 operations
- For each Enrolment entry with Code 'DBS' we find the corresponding Student from the ID
- $200 \times 15 = 3,000$  operations to do *both* the product and the selection.